

When Training Data Are Gone: Selective Integration of Pretrained Intrusion Detectors

Abstract—Effective intrusion detection benefits from diverse attack data, yet raw historical source records may no longer be accessible due to privacy and retention constraints, preventing further collaborative training. In this paper, we study selective integration of authorized pretrained detectors when the raw source records are unavailable but pretrained detectors and coarse source class counts remain available for reuse. Under a hard deployment budget, the goal is to choose a fixed detector set that preserves diverse source class coverage while maintaining predictive quality in the deployment domain. We present SILOSTITCH, which gives weighted source class coverage strict priority over deployment-domain predictive behavior rather than combining them into a single score. SILOSTITCH aligns heterogeneous detector outputs, uses held-out deployment records to assess deployment-domain predictive behavior, exactly maximizes weighted source class coverage under a logical byte budget via dynamic programming, and then refines the selected set without reducing coverage before fitting a shared stacker. To limit privacy leakage, we also protect uploaded score blocks via local differential privacy. Theoretically, we prove that calibration remains valid after adaptive selection and that the primary coverage problem is solved exactly, while perturbed scores satisfy label-conditional local differential privacy. Across four intrusion detection datasets with more than 60 candidate detectors per task, SILOSTITCH reduces communication by approximately half under a 50% selection-cost budget while preserving maximum weighted source class coverage and near-full-pool detection quality.

I. INTRODUCTION

Effective intrusion detection relies on training data that capture diverse traffic conditions and attack behaviors. However, data collected by a single organization typically reflect only its own operational environment, limiting the diversity of attack patterns available for training. Collaborative learning, particularly Federated Learning (FL), can mitigate such data silos by learning from distributed records without pooling them centrally [1]–[4]. Yet these approaches still require contributors to retain local training records and participate in additional model optimization. In practice, raw historical network logs may become inaccessible because they can contain sensitive operational and personally identifiable information [5], [6], and may need to be deleted under privacy and retention requirements such as the GDPR [7]. When these raw records are unavailable but authorized pretrained detectors and coarse source class counts remain, reusing the detectors offers a way to exploit complementary source knowledge without revisiting the original data.

Pretrained-model reuse can therefore operate after source records are no longer available [8]. However, existing approaches either select or aggregate pretrained models without explicitly accounting for heterogeneous deployment costs [9],

[10], or perform cost-aware detector selection rather than choosing one fixed detector set for deployment [11]. This gap motivates our research question: *How can a coordinator selectively integrate a fixed set of pretrained intrusion detectors under a deployment budget while preserving source intrusion-class coverage and maintaining detection performance in the deployment domain?*

Answering this question presents three challenges. First, detectors contributed by different parties may use different label spaces and support different target classes, making their outputs difficult to compare and integrate directly. Second, source class coverage and deployment-domain detection performance may favor different detectors: a detector with unique class coverage may no longer perform well after distribution shifts, whereas a detector that performs well in the deployment domain may add little new coverage. The coordinator should therefore preserve unique source coverage without filling the limited budget with detectors that are weak or redundant in the deployment domain. Third, detector selection must satisfy the deployment budget while remaining statistically reliable and privacy-aware. Because calibration data are used to choose detectors, their performance estimates must remain valid after selection, while uploaded prediction scores may still reveal sensitive information.

To address these challenges, we present SILOSTITCH, a framework for selective integration of pretrained intrusion detectors under deployment budgets. The key idea is to prioritize source class coverage over deployment-domain predictive performance: SILOSTITCH first maximizes weighted source class coverage under the deployment budget and then improves deployment-domain predictive behavior without reducing that coverage. This priority separates historical source coverage from deployment-domain predictive quality: coarse source class counts indicate which target classes were represented during training, whereas held-out deployment records indicate how well the pretrained detectors perform after distribution shifts [12]. SILOSTITCH first aligns heterogeneous detector outputs to a common label space and uses held-out deployment records to assess their deployment-domain predictive behavior. It then exactly solves the primary coverage problem through dynamic programming, refines the selected set without sacrificing coverage, and integrates the selected outputs with a shared stacker [13].

We further theoretically prove that calibration remains valid after adaptive detector selection and that the primary coverage problem is solved exactly. Because avoiding raw-data sharing alone does not eliminate privacy risks from model

outputs [14], SILOSTITCH also supports optional Laplace perturbation before score upload. The perturbed scores satisfy record-level, label-conditional local differential privacy; this guarantee applies to the uploaded scores, not to public labels or contributed detector parameters.

We evaluate SILOSTITCH on four intrusion detection datasets with over 60 candidate detectors per task. At a 50% selection-cost budget, SILOSTITCH maximizes weighted source class coverage while reducing total logical communication by 49.3% to 50.4%, demonstrating that diverse source coverage can be retained at roughly half the communication cost. Importantly, these savings incur only a small loss in detection quality: across all four datasets, the macro F1 gap from the full detector pool is limited to 0.028. Consistent with the exactness guarantee, the primary solver attains the same coverage as exhaustive search in all comparison instances. In summary, we make four contributions.

- We formulate budget-constrained selective integration of pretrained intrusion detectors when raw historical source records are unavailable and introduce an ordered objective that prioritizes source class coverage over deployment-domain predictive behavior.
- We develop SILOSTITCH to align heterogeneous detectors, exactly maximize weighted source class coverage under the budget, refine the selected set without reducing coverage, and integrate their outputs with a shared stacker.
- We prove that the calibration bounds remain valid after adaptive detector selection, that the primary coverage problem is solved exactly, and that local refinement preserves this optimum; we also establish label-conditional local differential privacy for randomized score uploads.
- We evaluate SILOSTITCH on four intrusion detection tasks, showing that a 50% selection-cost budget roughly halves communication while preserving maximum source class coverage and near-full-pool detection quality.

II. BACKGROUND AND RELATED WORK

We first review existing collaborative training and model integration methods, and then introduce the notion of differential privacy.

Retraining and Adaptation of Intrusion Detectors. When participants retain local training data and can rerun optimization, federated learning coordinates their parameter updates without transferring raw data [1]–[3]. Within intrusion detection, federated methods apply collaborative training across distributed network environments [6], [18]. Recent NIDS research also redesigns the training pipeline to withstand adverse data conditions. For example, RAPIER handles low-quality labels for encrypted malicious traffic [19], Rosetta uses TCP-aware augmentation to improve robustness across network environments [20], and CoLD denoises labels before fitting a classifier [21]. These approaches improve training or adaptation when the underlying records remain available to the learning procedure. In contrast, we consider the setting where the raw historical records are unavailable, while the pretrained

detectors and coarse class-count metadata remain authorized for reuse.

Pretrained Model Integration. When source training data are no longer available but their trained models remain accessible, pretrained model integration is the closest line of work. ZooD ranks a heterogeneous model zoo by cross-domain transferability and aggregates its top K models [9], while Model-GLUE studies selection and aggregation strategies for heterogeneous pretrained model zoos [10]. However, these methods focus on transfer or integration utility rather than the hard aggregate selection-cost constraint studied here.

Output-Level Ensembles and Cost-Aware Selection. Once detector outputs are available, stacked generalization can train a downstream classifier over them [13]. SIRAJ aggregates outputs from heterogeneous threat-intelligence sources using a pretrain-and-fine-tune framework [15], while FedMStacking uses stacking to support heterogeneous local models with misaligned labels [17]. To account for inference cost, SPIREL learns a per-item policy that trades off model invocations against classification utility [11]. These approaches cover complementary pieces of our setting: output ensembles combine available predictions, while SPIREL controls cost through per-item invocation decisions. SILOSTITCH instead makes a one-time subset decision after deployment-domain screening under a hard aggregate budget and then fits a shared classifier over the selected outputs.

Table I summarizes the combination of properties most relevant to our deployment setting: deployment-domain screening, fixed subset selection under a hard aggregate budget, source-class priority, and shared stacking over only the selected outputs. This combination motivates the ordered selection problem formalized next.

Local Differential Privacy. Prediction outputs can reveal information about their inputs, so a privacy guarantee must specify both where randomization occurs and which fields it protects. For example, PrivateKT perturbs locally produced predictions before uploading them for knowledge transfer [16], whereas differentially private federated trees protect the statistics exchanged while training a new model [22]. Our setting differs because SILOSTITCH uploads scores from selected detectors to fit a shared classifier while treating labels as public. We therefore specialize the local model of Duchi et al. [23] to this release.

Definition II.1 (Label-conditional local differential privacy). For a fixed public label y , a contributor applies a randomized mechanism $\mathcal{M}_y$ to the feature vector x . The upload $\mathcal{R}(x, y) = (\mathcal{M}_y(x), y)$ satisfies label-conditional $(\epsilon, 0)$ -local differential privacy if, for every fixed y , every pair x, x' , and every measurable output set $\mathcal{O}$,

$$\Pr[\mathcal{M}_y(x) \in \mathcal{O}] \leq e^\epsilon \Pr[\mathcal{M}_y(x') \in \mathcal{O}]. \quad (1)$$

This definition compares records with the same public label and protects the feature-derived portion of the upload, and the label itself remains public.

Table I
COMPARISON WITH RELATED METHODS. ✓ DENOTES DIRECT SUPPORT, ● RELATED SUPPORT, AND ○ NO SUPPORT.

Method	Model Integration			Subset Selection				Privacy	
	Pretrained Models	Heterogeneous	Score Stacking	Deployment Screening	Fixed Set	Hard Budget	Class Priority	Score	LDP
SIRAJ [15]	○	✓	●	○	○	○	○	○	○
PrivateKT [16]	○	○	●	○	○	○	○	✓	○
SPIREL [11]	✓	✓	●	○	●	●	○	○	○
ZooD [9]	✓	✓	●	○	✓	●	○	○	○
FedMStacking [17]	○	✓	✓	○	○	○	○	○	○
SILOSTITCH	✓	✓	✓	✓	✓	✓	✓	✓	✓

To instantiate Definition II.1, let a deterministic score map ψ have replacement sensitivity

$$\Delta_1(\psi) = \sup_{x, x'} \|\psi(x) - \psi(x')\|_1. \quad (2)$$

Adding independent noise $\eta_k \sim \text{Lap}(0, \Delta_1(\psi)/\epsilon)$ to every coordinate yields a mechanism that satisfies Definition II.1 [24]. If one record releases detector score vectors with parameters $\epsilon_1, \dots, \epsilon_s$, sequential composition gives the parameter $\sum_{i=1}^s \epsilon_i$. Under one record replacement, releases derived from unchanged records do not increase this privacy parameter. In SILOSTITCH, this mechanism applies to the selected detector score vectors before they are assembled into uploaded score blocks. Theorem V.1 establishes label-conditional local privacy for the resulting score upload under the public label adjacency relation.

III. PROBLEM STATEMENT

In this section, we formalize the problem of selective integration of pretrained intrusion detectors when the raw historical source records used to train them are unavailable. We first define the information available for detector reuse, map heterogeneous detector outputs to a common deployment label space, and model the communication cost of deploying a selected subset. We then specify the threat model for score uploads and formulate a budget-constrained selection objective that prioritizes source class coverage before deployment-domain predictive behavior.

A. System Model

We consider J contributors that provide a pool of M pretrained intrusion detectors for integration. The raw records originally used to fit these detectors are no longer available to the integration process, while the metadata needed to interpret and deploy the detectors remain available. For candidate i , let f_i denote the pretrained detector, with native probability output $p_i(x) \in \mathbb{R}_{\geq 0}^{h_i}$ satisfying $\mathbf{1}^\top p_i(x) = 1$, serialized model size K_i , and coarse source class-count vector $H_i = (H_{i,1}, \dots, H_{i,C})$ after semantic mapping, where $H_{i,c}$ records the source training count mapped to target class c . Because contributors may use different class definitions, we also associate each candidate with a semantic map that specifies how its native outputs correspond to the deployment label space.

We first make these heterogeneous outputs comparable. Let the deployment label space contain C target classes. For candidate i , the coordinator uses a semantic map

$$A_i \in \{0, 1\}^{(C+1) \times h_i}, \quad \mathbf{1}_{C+1}^\top A_i = \mathbf{1}_{h_i}^\top, \quad (3)$$

so each native output coordinate maps to exactly one aligned coordinate, while multiple native coordinates may map to the same target class. The coordinator then computes the aligned probability vector

$$q_i(x) = A_i p_i(x) \in \Delta^C, \quad (4)$$

where Δ^C is the probability simplex with $C + 1$ coordinates. The first C coordinates correspond to the predefined target classes. We reserve the final coordinate, denoted by $\perp$, for probability mass that cannot be mapped to any target class. By keeping this unmapped mass explicitly, later evaluation penalizes it instead of silently dropping it. We refer to $q_i(x)$ as candidate i 's aligned score vector for record x . For a record set $X = \{x_t\}_{t=1}^{|X|}$, we stack these vectors as

$$Q_i(X) = [q_i(x_t)^\top]_{t=1}^{|X|} \in \mathbb{R}^{|X| \times (C+1)}, \quad (5)$$

which we call the aligned score block of candidate i on X .

After making the detector outputs comparable, we use held-out deployment records to determine whether each pretrained detector remains useful under the deployment distribution. Before selection, the coordinator therefore holds a labeled calibration set $\mathcal{D}^{\text{cal}}$ from the deployment domain. These records are used only to evaluate the pretrained candidate pool and are disjoint from the records later used to fit the shared stacker and from the sealed test set. Every detector, semantic map, and source class-count vector is fixed before calibration begins. The test labels remain hidden until both detector selection and stacker fitting are complete, so they cannot influence detector selection or stacker fitting.

We next model the selection cost of choosing a detector. A selected detector must be delivered for use, and its aligned score block must later be uploaded for shared stacker fitting. We combine these two communication terms into a single selection cost measured in logical bytes. Here, logical bytes count application-level payload under the modeled workflow and exclude transport headers and implementation-specific protocol overhead. Let B_t denote the total number of stacker-fitting records across all contributors, and let d denote the

logical byte width assigned to one uploaded scalar, including either a probability value or a public label. For a selected set S , we define

$$C_{\text{sel}}(S) = \sum_{i \in S} b_i, \quad b_i = JK_i + dB_t(C+1). \quad (6)$$

The first term in b_i accounts for delivering candidate i to the J contributors, while the second accounts for uploading the aligned score block over all stacker-fitting records.

Given a budget fraction $\beta \in (0, 1]$, we set the hard selection-cost budget to

$$B = \left\lfloor \beta \sum_{i=1}^M b_i \right\rfloor. \quad (7)$$

The reference cost includes the entire candidate pool, including candidates that may later fail calibration. We additionally count the initial candidate upload and the stacker-label upload:

$$C_{\text{total}}(S) = \sum_{i=1}^M K_i + C_{\text{sel}}(S) + dB_t. \quad (8)$$

We use C_{sel} to enforce the selection-cost budget and C_{total} to compare the overall logical communication of SILOSTITCH with that of the full detector pool. This accounting covers integration and stacker fitting; distribution of the fitted stacker and subsequent inference traffic are outside the modeled communication cost.

B. Threat Model

We consider the privacy risk introduced when contributors upload detector scores for shared stacker fitting. Although raw historical source records are unavailable and never transferred during integration, the prediction scores computed from stacker-fitting records may still reveal information about those records. We therefore consider an honest-but-curious coordinator that follows the prescribed integration protocol but may attempt to infer record-level information from the uploaded scores.

We assume that each contributed detector, its semantic map, and its coarse source class counts are fixed and honestly reported before selection. Record labels used for stacker fitting are treated as public. When score protection is enabled, each contributor randomizes its selected score blocks locally before transmission, so the privacy guarantee does not rely on the coordinator performing the randomization. Our guarantee protects the feature-dependent information contained in these randomized score uploads; it does not protect the public labels, contributed detector parameters, or source class-count metadata. Malicious contributors that poison detector parameters, forge semantic maps or source class counts, or deviate from the prescribed randomization mechanism are outside our scope, as we focus on record-level privacy of score uploads.

C. Detector Selection Objective

The aligned outputs let us compare detectors in a common label space, but the budget may still prevent us from deploying all of them. We therefore need to decide which subset

to keep. Two considerations matter: we want the selected set to preserve source classes represented across the source domains, and we also want its detectors to remain useful in the deployment domain. These goals can favor different candidates. For example, one detector may be the only affordable model that covers a rare source class but have only moderate deployment-domain performance, while another may exhibit strong deployment-domain performance but add no new source-class coverage.

We handle this conflict with a strict priority: weighted source class coverage is maximized first under the selection-cost budget, and deployment-domain predictive behavior is compared only after that maximum is fixed. This is a lexicographic objective: the secondary objective can improve the choice among equally well-covered sets, but it can never compensate for losing primary coverage. We express this priority as

$$\text{lexmax}_{\emptyset \neq S \subseteq \hat{\mathcal{R}}} (C_{\text{src}}(S), \Phi(S)) \quad \text{s.t.} \quad \sum_{i \in S} b_i \leq B. \quad (9)$$

Here, $\hat{\mathcal{R}}$ is the pool that passes calibration, C_{src} is the weighted source class coverage defined later, and Φ summarizes deployment-domain predictive behavior. We solve the maximum source-coverage problem exactly and then apply coverage-preserving local refinement to improve deployment-domain predictive behavior. Equation 9 specifies this priority: SILOSTITCH solves the primary component exactly and refines the secondary component locally. Section V characterizes this refinement and explains why a single fixed weighted sum cannot enforce the same priority across arbitrary candidate pools.

IV. DESIGN OF SILOSTITCH

A. Overview of SILOSTITCH

We now describe how SILOSTITCH realizes the preceding problem formulation. Figure 1 summarizes the workflow in four steps. In Step 1, we use the semantic map A_i defined in the system model to transform each detector's native probabilities $p_i(x)$ into the aligned representation $q_i(x)$, so that heterogeneous detector outputs can be compared in a common label space. We also retain the source class-count metadata and selection cost introduced earlier for subsequent selection. These source counts are later converted into the coverage information used by the primary objective.

With the detector outputs made comparable, Step 2 uses held-out calibration records to evaluate how the pretrained detectors perform in the deployment domain. We use an upper confidence bound on Brier risk to determine whether each detector enters the calibration-eligible pool, while retaining additional classwise utility and prediction-pattern statistics for later refinement. Based on this eligible pool, Step 3 first uses the source-support information and detector selection costs to find the maximum weighted source class coverage allowed by the selection-cost budget. We then use the retained calibration statistics to improve deployment-domain predictive behavior without reducing the achieved coverage. Finally, once the

Algorithm 1 SILOSTITCH Procedure

Require: $\{f_i, A_i, H_i, K_i\}_{i=1}^M, \mathcal{D}^{\text{cal}}, \{\mathcal{D}_j^{\text{stk}}\}_{j=1}^J, B, \epsilon_b \in (0, \infty]$

Ensure: $\hat{S}, \hat{f}_{\hat{S}}$

- 1: Construct (Γ, w, b) from the fixed source metadata and cost model
 - 2: Evaluate candidates on $\mathcal{D}^{\text{cal}}$ to obtain $\hat{\mathcal{R}}$ and the statistics used by Φ
 - 3: Construct Φ from the retained calibration statistics
 - 4: $(W^*, S_0) \leftarrow \text{COVER}(\hat{\mathcal{R}}, \Gamma, w, b, B)$
 - 5: $\hat{S} \leftarrow \text{REFINE}(S_0, W^*, \Phi, B)$
 - 6: Coordinator $\rightarrow$ Contributors: $\{f_i\}_{i \in \hat{S}}$
 - 7: **for** $j = 1, \dots, J$ **in parallel do**
 - 8: $Z_j \leftarrow \text{concat}_{i \in \hat{S}} Q_i(X_j^{\text{stk}})$
 - 9: **if** $\epsilon_b < \infty$ **then**
 - 10: $(\Xi_j)_{t,i,k} \stackrel{\text{iid}}{\sim} \text{Lap}(0, 2/\epsilon_b)$
 - 11: $\tilde{Z}_j \leftarrow Z_j + \Xi_j$
 - 12: **else**
 - 13: $\tilde{Z}_j \leftarrow Z_j$
 - 14: **end if**
 - 15: Contributor $j \rightarrow$ Coordinator: $(\tilde{Z}_j, Y_j^{\text{stk}})$
 - 16: **end for**
 - 17: $\hat{f}_{\hat{S}} \leftarrow \text{FIT}(\bigcup_j (\tilde{Z}_j, Y_j^{\text{stk}}))$
 - 18: **return** $(\hat{S}, \hat{f}_{\hat{S}})$
-

selected detector set is fixed, Step 4 uses only its aligned score blocks to fit one shared stacker and keeps the same selected layout for subsequent sealed-test inference. The following subsections define the operations in these four steps in detail.

B. Selection and Stacking Workflow

Algorithm 1 summarizes the four steps above. Let j index contributors and i index candidate detectors, and write $\mathcal{D}_j^{\text{stk}} = (X_j^{\text{stk}}, Y_j^{\text{stk}})$ for contributor j 's stacker-fitting records and labels. The calibration partition is first used to evaluate the pretrained candidates, form the eligible pool, and compute the exact primary coverage value. We then refine one exact primary solution using the retained calibration evidence. Only after the selected set $\hat{S}$ is fixed do we use the independent stacker-fitting records to train the shared stacker. This separation prevents the records used to fit the stacker from also influencing which detector outputs become its input features.

The contributors then evaluate only the selected detectors on their stacker-fitting records. If score protection is enabled, they randomize the selected aligned score blocks locally before upload; otherwise the same workflow is used without added noise.

C. Brier-Based Evaluation and Screening

After aligning every candidate to the same target label space through $q_i(x)$, we evaluate the aligned outputs on held-out deployment records to obtain comparable statistics for screening and refinement. For this assessment, we use a full-vector Brier score because it is bounded, evaluates the complete aligned probability vector, and naturally penalizes probability mass assigned to the unmapped $\perp$ coordinate.

For calibration record (x_t, y_t) , let $e_{y_t} \in \mathbb{R}^{C+1}$ be the one-hot label vector with a zero $\perp$ coordinate. For candidate i , we define

$$\ell_{i,t} = \frac{1}{2} \|q_i(x_t) - e_{y_t}\|_2^2, \quad g_{i,t} = 1 - \ell_{i,t} \in [0, 1]. \quad (10)$$

Thus, $g_{i,t}$ measures the quality of the whole aligned prediction rather than one coordinate at a time. If a detector places probability on an unmapped output, that probability remains in the vector and lowers the utility.

Eligibility screening. We first determine whether each candidate detector has sufficiently low deployment-domain risk to remain eligible. Using only its empirical mean can be misleading when the calibration set is finite, so we screen with an upper confidence bound on deployment-domain Brier risk. With N calibration records and failure probability $\delta \in (0, 1)$, we compute

$$\hat{\mu}_i = \frac{1}{N} \sum_{t=1}^N g_{i,t}, \quad r = \sqrt{\frac{\log(4M/\delta)}{2N}}, \quad (11)$$

and define

$$R_i^+ = 1 - \hat{\mu}_i + r. \quad (12)$$

We keep candidate i only if $R_i^+ \leq \tau_L$, giving the calibration-eligible pool

$$\hat{\mathcal{R}} = \{i : R_i^+ \leq \tau_L\}. \quad (13)$$

In other words, eligibility depends on both the observed deployment-domain performance and the uncertainty caused by finite calibration data.

Classwise scores. An acceptable average risk does not guarantee that a detector remains useful for every target class. We therefore keep a conservative classwise utility estimate for each detector-class pair and use these scores later when we refine sets that already achieve maximum source class coverage. For class c , let

$$N_c = \sum_t \mathbf{1}\{y_t = c\}, \quad (14)$$

and define

$$\hat{\mu}_{i,c} = \frac{1}{N_c} \sum_{t: y_t=c} g_{i,t}. \quad (15)$$

We require every calibration class to appear at least once, so $N_c > 0$. We then compute

$$r_c = \sqrt{\frac{\log(4MC/\delta)}{2N_c}}, \quad a_{i,c} = [\hat{\mu}_{i,c} - r_c]_+, \quad (16)$$

where $[z]_+ = \max\{0, z\}$. The value $a_{i,c}$ is a lower confidence bound on the deployment-domain utility of candidate i on class c .

Prediction similarity. Two detectors can have similar average utility and still make nearly identical predictions. Selecting both may then consume budget without adding much new predictive information. To capture this redundancy, we compare




Figure 1. Overview of SILOSTITCH. Step 1 maps the heterogeneous outputs of the pretrained detector pool to a common $C + 1$ -coordinate label space and retains the metadata needed for selection. Step 2 evaluates the aligned detectors on held-out calibration records and forms the calibration-eligible pool. Step 3 exactly maximizes weighted source class coverage under the selection-cost budget and then refines deployment-domain predictive behavior without reducing the achieved coverage. Step 4 fits a shared stacker from the selected aligned outputs and reuses the same selected layout for sealed-test inference.

the aligned outputs of the candidates on a fixed class-stratified set of calibration indices $\mathcal{A} \subseteq \{1, \dots, N\}$. We form

$$z_i = \text{vec}([q_i(x_t)]_{t \in \mathcal{A}}), \quad \kappa_{i,k} = \left[\frac{\langle z_i, z_k \rangle}{\|z_i\|_2 \|z_k\|_2} \right]_0^1. \quad (17)$$

The clipping only removes numerical roundoff outside $[0, 1]$. Later, we use these similarities in a representativeness term: rather than rewarding pairwise difference for its own sake, we ask whether the selected set contains detectors whose prediction patterns represent the wider eligible pool.

D. Coverage-Prioritized Detector Selection

The calibration statistics tell us how the pretrained detectors perform in the deployment domain, but they do not tell us which historical source classes would disappear if particular detectors were removed. That information comes from the coarse source class counts retained with the candidate models. We therefore use source class coverage as the primary objective and use the calibration statistics only after the best achievable coverage has been fixed.

Weighted source class coverage. Recall that $H_{i,c}$ denotes candidate i 's mapped source count for target class c . A very small count may provide weak evidence that a class was meaningfully represented during training, so we first convert the count into a support indicator. For a minimum support threshold $n_{\min} > 0$, define

$$\Gamma_{i,c} = \mathbf{1}\{H_{i,c} \geq n_{\min}\}. \quad (18)$$

Thus, $\Gamma_{i,c} = 1$ indicates that candidate i has sufficient source support for class c to count toward the primary objective. Let $\Gamma_i = (\Gamma_{i,1}, \dots, \Gamma_{i,C})$ denote candidate i 's source-support vector.

We also give more weight to source classes that are scarce across the whole candidate pool. Let $H_c = \sum_i H_{i,c}$ denote the aggregate source count of class c , and let $\mathcal{C}_+ = \{c : H_c > 0\}$

denote the source-supported classes. We assume $\mathcal{C}_+ \neq \emptyset$ and set

$$w_c = \begin{cases} \frac{H_c^{-1/2}}{\sum_{c' \in \mathcal{C}_+} H_{c'}^{-1/2}}, & c \in \mathcal{C}_+, \\ 0, & c \notin \mathcal{C}_+, \end{cases} \quad \sum_{c=1}^C w_c = 1. \quad (19)$$

The inverse-square-root rule gives greater weight to less represented source-supported classes while moderating frequency differences relative to inverse-frequency weighting.

For selected set S , we define weighted source class coverage as

$$C_{\text{src}}(S) = \sum_{c=1}^C w_c \mathbf{1}\{\exists i \in S : \Gamma_{i,c} = 1\}. \quad (20)$$

A class contributes once as soon as at least one selected detector provides sufficient source support for it; selecting more detectors for the same class does not increase the primary value. Therefore, C_{src} measures how much historical source class coverage is retained by the selected set, not how accurately that set predicts deployment records.

Secondary predictive objective. After the primary coverage value is fixed, several feasible detector sets may still attain exactly the same coverage. We compare such sets using three secondary objective components that capture different aspects of deployment-domain prediction quality: record-level utility, classwise utility, and prediction-pattern representativeness. With $h_\tau(z) = 1 - e^{-z/\tau}$, define

$$\begin{aligned} U(S) &= \frac{1}{N} \sum_{t=1}^N h_{\tau_U} \left(\sum_{i \in S} g_{i,t} \right), \\ C_{\text{pred}}(S) &= \frac{1}{C} \sum_{c=1}^C h_{\tau_C} \left(\sum_{i \in S} a_{i,c} \right), \\ D_{\text{rep}}(S) &= \frac{1}{|\widehat{\mathcal{R}}|} \sum_{k \in \widehat{\mathcal{R}}} \max_{i \in S} \kappa_{i,k}. \end{aligned} \quad (21)$$

The term $U(S)$ rewards sets whose members perform well across the calibration records. The term $C_{\text{pred}}(S)$ keeps attention on classwise utility by using the conservative scores $a_{i,c}$. The term $D_{\text{rep}}(S)$ rewards sets whose prediction patterns represent those observed across the eligible pool.

The saturating functions in U and C_{pred} impose diminishing returns. Once several selected detectors already perform well on the same record or class, another similar detector contributes less. This makes it harder for redundant prediction patterns to dominate the secondary comparison. We define $\max_{i \in \emptyset} \kappa_{i,k} = 0$ and combine the three terms as

$$\Phi(S) = \lambda_U U(S) + \lambda_C C_{\text{pred}}(S) + \lambda_D D_{\text{rep}}(S), \quad (22)$$

where $\tau_U, \tau_C > 0$, $\lambda_U, \lambda_C, \lambda_D \geq 0$, and $\lambda_U + \lambda_C + \lambda_D = 1$. We use Φ only after the primary coverage value has been fixed, so no gain in deployment-domain predictive behavior can compensate for a loss in source class coverage.

Budgeted selection procedure. We now solve the primary problem: among the calibration-eligible candidates, what is the largest weighted source class coverage that fits within the selection-cost budget? Define

$$\mathcal{F}_B = \left\{ \emptyset \neq S \subseteq \hat{\mathcal{R}} : \sum_{i \in S} b_i \leq B \right\}, \quad (23)$$

$$W^* = \max_{S \in \mathcal{F}_B} C_{\text{src}}(S).$$

We assume $\mathcal{F}_B \neq \emptyset$; otherwise, the budget admits no eligible detector. Because C_{src} depends only on which classes are covered, we encode every source-support vector Γ_i as a C -bit mask and run a sparse dynamic program over reachable masks. For each mask, we retain only the least-cost detector set that produces it. Any more expensive set with the same mask is dominated because every future detector changes the two masks in exactly the same way.

The dynamic program returns both the exact optimum W^* and a least-cost set S_0 that attains it. Starting from this least-cost exact solution leaves as much unused budget as possible for improving deployment-domain predictive behavior. With a tolerance $\epsilon_{\text{loc}} \geq 0$, we first consider additional candidates according to their secondary gain per byte and accept an addition only when its gain exceeds ϵ_{loc} . We then examine one-for-one replacements and accept a swap only when it remains budget-feasible, keeps coverage equal to W^* , and improves Φ by more than the same tolerance. Every accepted refinement therefore improves deployment-domain predictive behavior without weakening the exact primary guarantee.

E. Shared Stacker Learning

Once selection fixes $\hat{S}$, only the aligned outputs of the selected detectors enter the shared stacker. We use the same fixed detector order for stacker fitting and inference. For input x , we define the selected representation

$$\phi_{\hat{S}}(x) = \text{concat}_{i \in \hat{S}} q_i(x) \in \mathbb{R}^{|\hat{S}|(C+1)}. \quad (24)$$

The selected detector set determines which score blocks are used, while this fixed order determines their feature-column positions.

Local score perturbation. Although raw historical source records are unavailable, the scores computed from stacker-fitting records may still reveal information about those records. When score protection is enabled, each contributor therefore randomizes every coordinate of each selected detector score vector before upload:

$$\tilde{q}_i(x_t) = q_i(x_t) + \eta_{i,t}, \quad \eta_{i,t,k} \stackrel{\text{iid}}{\sim} \text{Lap}\left(0, \frac{2}{\epsilon_b}\right). \quad (25)$$

The resulting stacker input is

$$\tilde{\phi}_{\hat{S}}(x) = \text{concat}_{i \in \hat{S}} \tilde{q}_i(x). \quad (26)$$

Section V later proves the corresponding label-conditional local privacy guarantee and its composition across selected detector score vectors.

Shared stacker fitting. Let $\mathcal{D}^{\text{stk}} = \biguplus_{j=1}^J \mathcal{D}_j^{\text{stk}}$ denote the pooled stacker-fitting records. After selection is complete, this independent partition supplies the labels and selected score features used to train one shared classifier:

$$\hat{f}_{\hat{S}} = \text{Fit}_{\text{global}} \left(\{(\tilde{\phi}_{\hat{S}}(x_t), y_t) : (x_t, y_t) \in \mathcal{D}^{\text{stk}}\} \right). \quad (27)$$

For the nonprivate configuration, we simply set $\tilde{\phi}_{\hat{S}} = \phi_{\hat{S}}$. Because both configurations use only the columns indexed by $\hat{S}$, unselected detectors cannot influence the learned stacker.

In a distributed realization of the stacker-fitting step, contributors generate the Laplace noise locally before transmission. Our centralized privacy experiment samples the same distribution only to isolate the utility effect of score randomization. For sealed-test inference, we keep the selected detector set and aligned feature layout fixed and reuse the fitted shared stacker without retraining any contributed detector.

V. THEORETICAL ANALYSIS

We now analyze the guarantees that support the workflow above. First, we show what the local score randomization protects. Second, we show that the calibration bounds remain valid even after we use those same calibration statistics to choose detectors. Third, we prove that the primary coverage solver is exact and characterize what the local secondary refinement guarantees. Finally, we examine how a change in the contributor mixture can affect the interpretation of the selected detectors after deployment.

A. Local Privacy of Score Uploads

Recall that the shared stacker is trained from selected detector scores rather than from the historical source records. Because these scores are still record dependent, we ask what information the contributor-side Laplace mechanism protects. We use the label-conditional adjacency relation defined earlier: the record label is public, and two neighboring records (x, y) and (x', y) differ only in their feature vectors.

Before any stacker-fitting record is processed, we fix the selected detector set and all semantic maps. For a batched

upload, neighboring batches have the same size, the same row alignment, and the same sequence of public labels, but they differ in one feature vector. Each contributor samples fresh noise independently across records, selected detector score vectors, and output coordinates. Under this execution model, we obtain the following guarantee.

Theorem V.1 (Label-conditional local score privacy). Suppose that the selected detectors and their semantic maps are fixed and that $0 < \epsilon_b < \infty$, and suppose that each contributor uses fresh private randomness in Eq. 25. Then one released detector score vector satisfies $(\epsilon_b, 0)$ -local differential privacy under the preceding label-conditional adjacency relation. If $s = |\hat{S}|$ detector score vectors are released for the same record, their joint release satisfies $(s\epsilon_b, 0)$ -local differential privacy. Under one-record replacement, releasing the entire contributor batch has the same record-level guarantee. More generally, assigning parameter ϵ_i to detector score vector i gives $(\sum_{i \in \hat{S}} \epsilon_i, 0)$ -local differential privacy.

Proof. Because $q_i(x)$ and $q_i(x')$ are probability vectors, their replacement sensitivity satisfies

$$\|q_i(x) - q_i(x')\|_1 \leq 2. \quad (28)$$

Therefore, adding coordinatewise Laplace noise with scale $2/\epsilon_b$ bounds the density ratio of one released detector score vector by e^{ϵ_b} . Sequential composition over the s selected detector score vectors gives the factor $e^{s\epsilon_b}$.

For a batched upload, neighboring batches differ in only one row. Since different rows use independent randomization, parallel composition across the unchanged rows does not increase the record-level privacy parameter. Finally, appending the unchanged public labels does not change the density ratio. $\square$

The theorem also clarifies how the privacy budget interacts with detector selection. If we want one total per-record score budget ϵ , we can assign $\epsilon_b = \epsilon/s$ to each of the s selected detector score vectors. The corresponding noise scale becomes $2s/\epsilon$. Thus, under a fixed total score-privacy budget, selecting fewer detector score vectors also reduces the noise required for each released vector. The guarantee applies only to the randomized score upload, and privacy loss composes again if the same record is released in repeated rounds.

B. Calibration under Heterogeneous Contributors

The privacy result concerns the stacker-fitting partition. Detector selection, in contrast, depends on the separate calibration partition introduced in the system model. We therefore ask a different question here: after we use the calibration statistics to choose a detector subset, can we still trust the confidence bounds computed before selection?

The calibration records may come from different contributors and classes, so we do not assume that every record follows one identical distribution. Instead, let j_t identify the

contributor associated with calibration record t , and condition on the observed contributor-label design

$$\mathcal{I} = \{(j_t, y_t)\}_{t=1}^N \quad (29)$$

and on the pretrained candidate pool $\mathcal{Q}$. Let X_t denote the random feature vector for calibration record t , and define

$$G_{i,t} = 1 - \frac{1}{2} \|q_i(X_t) - e_{y_t}\|_2^2 \in [0, 1], \quad (30)$$

whose observed realization is $g_{i,t}$ in Eq. 10. Conditioned on $(\mathcal{Q}, \mathcal{I})$, we assume the records are independent, while allowing their conditional distributions to differ across records. We define

$$\mu_i = \frac{1}{N} \sum_{t=1}^N \mathbb{E}[G_{i,t} \mid \mathcal{Q}, \mathcal{I}], \quad R_i = 1 - \mu_i, \quad (31)$$

and

$$\mu_{i,c} = \frac{1}{N_c} \sum_{t: y_t=c} \mathbb{E}[G_{i,t} \mid \mathcal{Q}, \mathcal{I}]. \quad (32)$$

These quantities average performance over the actual mix of calibration records observed by the coordinator; they are not equal-contributor averages or worst-contributor guarantees.

Theorem V.2 (Simultaneous calibration). Under the preceding assumptions, with probability at least $1 - \delta$, the following inequalities hold simultaneously for every candidate i and class c :

$$R_i \leq R_i^+ \leq R_i + 2r, \quad (33)$$

$$\max\{0, \mu_{i,c} - 2r_c\} \leq a_{i,c} \leq \mu_{i,c}. \quad (34)$$

Consequently, these inequalities remain valid for every candidate in any set selected adaptively from the calibration statistics.

Proof. Conditioned on $(\mathcal{Q}, \mathcal{I})$, Hoeffding's inequality gives

$$\Pr\{|\hat{\mu}_i - \mu_i| > r \mid \mathcal{Q}, \mathcal{I}\} \leq 2e^{-2Nr^2} = \frac{\delta}{2M}.$$

A union bound over the M candidates therefore uses at most $\delta/2$ failure probability. Likewise,

$$\Pr\{|\hat{\mu}_{i,c} - \mu_{i,c}| > r_c \mid \mathcal{Q}, \mathcal{I}\} \leq 2e^{-2N_c r_c^2} = \frac{\delta}{2MC},$$

so a second union bound over all MC candidate-class pairs uses the remaining $\delta/2$. Thus, with probability at least $1 - \delta$, all overall and classwise inequalities hold at the same time.

On this event, $1 - \hat{\mu}_i + r$ lies between R_i and $R_i + 2r$, while $\hat{\mu}_{i,c} - r_c$ lies between $\mu_{i,c} - 2r_c$ and $\mu_{i,c}$. Applying $[\cdot]_+$ gives the classwise result. The key point is that this one event covers all $M + MC$ quantities before selection occurs. We can therefore choose a data-dependent subset without paying another union bound over the 2^M possible detector sets. $\square$

The same simultaneous event also tells us how the eligibility screen can affect the best source coverage available to the selector. For any candidate pool $\mathcal{E}$, let $W(\mathcal{E})$ denote the largest

budget-feasible weighted source class coverage available from that pool, with value zero if no nonempty set is feasible. Define

$$\mathcal{R}^- = \{i : R_i \leq \tau_L - 2r\}, \quad \mathcal{R}^0 = \{i : R_i \leq \tau_L\}. \quad (35)$$

Then

$$\mathcal{R}^- \subseteq \widehat{\mathcal{R}} \subseteq \mathcal{R}^0, \quad W(\mathcal{R}^-) \leq W(\widehat{\mathcal{R}}) \leq W(\mathcal{R}^0). \quad (36)$$

The left inclusion follows from $R_i^+ \leq R_i + 2r$, while the right inclusion follows from $R_i \leq R_i^+$. Therefore, if one primary-optimal set over $\mathcal{R}^0$ already lies entirely in $\mathcal{R}^-$, calibration screening does not reduce the best achievable primary coverage. This upper-confidence gate therefore provides a finite-sample eligibility screen before budgeted selection.

C. Exact Primary Coverage under a Hard Budget

Once calibration fixes the eligible pool, the primary selection problem has a simple interpretation. Each detector has a selection cost and covers a subset of source classes, while the hard budget may prevent us from selecting all detectors at once. We first characterize the difficulty of this problem and then show why the bitmask dynamic program used by SILOSTITCH still returns the exact primary optimum for the processed instances.

Theorem V.3 (Primary hardness and exact parameterized solver). Computing W^* is NP-hard, even when every candidate is eligible, every cost is one, and all class weights are uniform. Nevertheless, the bitmask dynamic program returns a nonempty feasible set S_0 such that

$$C_{\text{src}}(S_0) = W^*, \quad (37)$$

with minimum total selection cost among the sets attaining W^* . It stores at most 2^C masks and examines at most $|\widehat{\mathcal{R}}|2^C$ state transitions. Thus, the primary stage is fixed-parameter tractable in the number of target classes C . In other words, the transition bound is linear in the number of eligible detectors but exponential in the number of target classes.

Proof. For hardness, take a Maximum Coverage instance with universe $\{1, \dots, C\}$, subsets $\mathcal{U}_1, \dots, \mathcal{U}_M$, and cardinality limit k . Set $\Gamma_{i,c} = 1$ exactly when $c \in \mathcal{U}_i$, and set $b_i = 1$, $B = k$, and $w_c = 1/C$. Maximizing C_{src} then solves that Maximum Coverage instance up to the constant factor $1/C$.

For exactness, process the eligible candidates one at a time. After the first t candidates, associate every reachable union mask h with the least-cost detector set that produces that mask. When we process the next candidate, every subset either excludes it and keeps an earlier state, or includes it and extends an earlier mask to $h \vee \Gamma_{t+1}$. If two sets produce the same mask, the cheaper one dominates the more expensive one because every future candidate changes the two masks in exactly the same way. By induction, the final table therefore contains the least-cost realization of every reachable mask. Scanning the masks whose costs do not exceed B returns exactly W^* ; among masks attaining this value, we choose the least-cost stored realization.

A C -bit mask has at most 2^C values, and every eligible candidate extends each reachable mask at most once. If all feasible masks have zero weighted coverage, we return the least-cost feasible singleton so that the downstream stacker still receives a nonempty input without changing the primary value. $\square$

The transition count does not include reconstructing the selected detector set or sorting candidates for the later refinement stage. In our implementation, we process at most 10^6 mask states. The exactness guarantee applies only when the full dynamic program completes without truncation at this implementation cap.

D. Secondary Objective and Local Refinement

The exact primary solver fixes the best source coverage that the budget can support, but many detector sets may still attain that same value. We use the secondary objective Φ to choose among them. Before stating the local-refinement guarantee, we first explain why we do not replace the ordered objective with one weighted sum.

For a fixed problem instance, define the set of achievable primary values as

$$\mathcal{V} = \{C_{\text{src}}(S) : S \in \mathcal{F}_B\}. \quad (38)$$

When $\mathcal{V}$ contains at least two values, let Δ_{inst} be the smallest positive gap between them. Because $0 \leq \Phi(S) \leq 1$, any coefficient $K > 1/\Delta_{\text{inst}}$ makes every maximizer of $K C_{\text{src}} + \Phi$ respect the lexicographic order for that one fixed instance. However, Δ_{inst} depends on the class weights and candidate pool and can become arbitrarily small. Therefore, no universal finite K enforces the required priority across arbitrary deployments. The ordered formulation in Eq. 9 states the priority directly and avoids instance-specific scale tuning.

We next ask whether the secondary score has a useful diminishing-returns property: once a selected set already captures a prediction pattern, adding another similar detector should help less. This standard set-function property is called submodularity.

Proposition V.4 (Secondary structure). For fixed calibration statistics, U , C_{pred} , and D_{rep} are normalized, monotone, submodular set functions. Consequently, their nonnegative weighted combination Φ has the same properties.

Proof. For U and C_{pred} , each summand applies the nondecreasing concave function h_τ to a nonnegative modular sum. Therefore, the marginal gain from adding the same candidate cannot increase as the selected set grows. For D_{rep} , eligible detector k contributes $\max_{i \in S} \kappa_{i,k}$, whose marginal gain from adding candidate a is

$$\max \left\{ 0, \kappa_{a,k} - \max_{i \in S} \kappa_{i,k} \right\}.$$

As S grows, the existing maximum can only increase, so this gain can only decrease. Averaging and taking a nonnegative weighted sum preserve normalization, monotonicity, and submodularity. $\square$

This submodularity does not by itself give a global approximation guarantee for the deployed secondary refinement. The family of sets that preserve exact primary coverage,

$$\mathcal{F}^* = \{S \in \mathcal{F}_B : C_{\text{src}}(S) = W^*\}, \quad (39)$$

does not have the simple structure required by standard submodular-knapsack results: removing a detector can destroy the required coverage, and an arbitrary exchange need not restore it. Therefore, familiar global guarantees such as $1 - 1/e$ do not directly apply after we require exact primary coverage.

For $S \in \mathcal{F}^*$, let $\mathcal{N}_1(S)$ contain every set obtained by removing one selected candidate and adding one unselected candidate while preserving budget feasibility and coverage W^* . We can then state a local certificate for a refinement that reaches a one-swap fixed point.

Theorem V.5 (Coverage preservation and one-swap certificate). Given an exact primary solution, suppose the replacement phase terminates after a complete pass in which no budget-feasible, coverage-preserving swap improves Φ by more than ε_{loc} . Then the refinement returns $\hat{S} \in \mathcal{F}^*$ such that

$$\Phi(S') \leq \Phi(\hat{S}) + \varepsilon_{\text{loc}}, \quad \forall S' \in \mathcal{N}_1(\hat{S}). \quad (40)$$

Proof. The initial set has coverage W^* . Since W^* is already the maximum feasible coverage and C_{src} is monotone, every feasible addition keeps that primary value. Every accepted replacement is also checked explicitly for budget feasibility and equality to W^* . Each accepted operation improves Φ by more than ε_{loc} . Because the candidate pool contains only finitely many subsets, any sequence of accepted operations that strictly increases Φ cannot revisit a set. Under the stated termination condition, no budget-feasible, coverage-preserving one-swap neighbor improves Φ by more than the tolerance, which yields Eq. 40. $\square$

The refinement examines additions and one-for-one replacements rather than all possible multi-detector exchanges. It always preserves the exact primary optimum. If the replacement phase reaches the termination condition in Theorem V.5, the returned set additionally satisfies the stated one-swap certificate.

E. Coverage Certificates under Contributor Mixture Shift

The source-coverage objective records which classes were represented in the historical source domains. It does not, by itself, guarantee that a detector supporting a class still performs well on that class after deployment. The held-out calibration records provide calibration evidence, but the contributor mixture observed during deployment may itself differ from the mixture seen during calibration. We therefore connect historical source support with classwise predictive utility and quantify how a change in contributor mixture can weaken that connection.

This analysis does not change which set SILOSTITCH selects. Instead, after selection, we use the already computed classwise calibration bounds to ask which historically covered

classes also have conservative calibration evidence of classwise utility. For any threshold vector $\mathbf{u} \in \mathbb{R}^C$, define

$$C_{\text{cert}}(S; \mathbf{u}) = \sum_{c=1}^C w_c \mathbf{1}\{\exists i \in S : \Gamma_{i,c} = 1 \wedge a_{i,c} \geq u_c\}. \quad (41)$$

A class contributes to this value only when at least one selected detector both reports enough historical source support and has a conservative calibration utility above the threshold; if $u_c > 1$, no detector can satisfy that classwise-utility condition.

To express the same notion using the underlying population utilities rather than their calibration lower bounds, let $C_{\mu}^a(S; \mathbf{u})$ replace the condition $a_{i,c} \geq u_c$ by $\mu_{i,c}^a \geq u_c$, where $a \in \{\text{cal}, \text{dep}\}$ denotes calibration and deployment. We consider a mixture-shift model in which each detector's mean Brier risk within a contributor-class cell remains stable while only the proportions of those cells change between calibration and deployment. This additional stability assumption is used only for the post-selection analysis in this subsection. Let $R_{i,j,c} \in [0, 1]$ be the mean Brier risk of detector i in contributor-class cell (j, c) . We take $\pi_{j,c}^{\text{cal}} = N^{-1} \sum_{t=1}^N \mathbf{1}\{j_t = j, y_t = c\}$ as the record-weighted contributor-class composition of the calibration set, and let $\pi_{j,c}^{\text{dep}}$ denote the corresponding deployment mixture. For $a \in \{\text{cal}, \text{dep}\}$, define

$$\pi_c^a = \sum_j \pi_{j,c}^a, \quad \pi_{j|c}^a = \frac{\pi_{j,c}^a}{\pi_c^a}, \quad (42)$$

assuming $\pi_c^a > 0$ for every class considered in the certificate. We then define

$$R_i^a = \sum_{j,c} \pi_{j,c}^a R_{i,j,c}, \quad \mu_{i,c}^a = \sum_j \pi_{j|c}^a (1 - R_{i,j,c}), \quad (43)$$

and let $\text{TV}(p, q) = \frac{1}{2} \|p - q\|_1$. Under the cellwise mean-risk assumption, R_i^{cal} and $\mu_{i,c}^{\text{cal}}$ coincide with the calibration quantities in Eqs. 31 and 32, respectively.

Because $1 - R_{i,j,c} \in [0, 1]$, a within-class bound $\text{TV}(\pi_{\cdot|c}^{\text{dep}}, \pi_{\cdot|c}^{\text{cal}}) \leq \eta_c$ implies $|\mu_{i,c}^{\text{dep}} - \mu_{i,c}^{\text{cal}}| \leq \eta_c$.

Theorem V.6 (Certified coverage under mixture shift). Let $\boldsymbol{\vartheta} \in [0, 1]^C$ be prespecified and suppose that $\text{TV}(\pi_{\cdot|c}^{\text{dep}}, \pi_{\cdot|c}^{\text{cal}}) \leq \eta_c$ for every class c , where $\eta_c \in [0, 1]$. On the simultaneous calibration event, every adaptively selected set $\hat{S}$ satisfies

$$\begin{aligned} C_{\mu}^{\text{cal}}(\hat{S}; \boldsymbol{\vartheta} + \boldsymbol{\eta} + 2\mathbf{r}) &\leq C_{\text{cert}}(\hat{S}; \boldsymbol{\vartheta} + \boldsymbol{\eta}) \\ &\leq C_{\mu}^{\text{dep}}(\hat{S}; \boldsymbol{\vartheta}) \leq C_{\text{src}}(\hat{S}), \end{aligned} \quad (44)$$

where $(\mathbf{r})_c = r_c$ and $(\boldsymbol{\eta})_c = \eta_c$. Moreover, if $\text{TV}(\pi_{\cdot|c}^{\text{dep}}, \pi_{\cdot|c}^{\text{cal}}) \leq \eta$ for some $\eta \in [0, 1]$, then

$$R_i^{\text{dep}} \leq R_i^{\text{cal}} + \eta. \quad (45)$$

Proof. The simultaneous calibration result and the mixture-shift assumption give

$$\mu_{i,c}^{\text{cal}} \geq \vartheta_c + \eta_c + 2r_c \implies a_{i,c} \geq \vartheta_c + \eta_c \implies \mu_{i,c}^{\text{dep}} \geq \vartheta_c.$$

Combining these implications with the fixed historical-support condition $\Gamma_{i,c} = 1$, taking the existence condition over the selected detectors, and summing the class weights yields Eq. 44. For the final result, total variation bounds the change in the expectation of any function taking values in $[0, 1]$. Applying this fact to the contributor–class risks gives Eq. 45. $\square$

The theorem separates three quantities that serve different purposes. C_{src} is the historical source-coverage value optimized during selection. C_{cert} keeps only the part of that coverage that also has conservative calibration evidence. C_{μ}^{dep} describes the corresponding deployment utility after mixture shift. Thus, SILOSTITCH guarantees the historical coverage objective exactly, while this separate post-selection analysis states when that historical coverage is also supported by deployment predictive evidence.

VI. EVALUATION

We evaluate whether selective integration of pretrained intrusion detectors can deliver the three properties targeted by our design: preserving source class coverage, maintaining useful deployment-domain detection quality, and reducing total logical communication. We also study when these benefits become weaker by varying source heterogeneity, secondary design choices, privacy strength, and problem scale. To keep the discussion organized, we answer five questions.

- **RQ1: Detection quality under a budget.** How much source coverage and deployment-domain detection quality do we retain when only part of the full-pool selection cost is available?
- **RQ2: Contributor label heterogeneity.** How does stronger heterogeneity across source contributors affect candidate construction, source coverage, and final detection quality?
- **RQ3: Selection components and learners.** How sensitive are the selected sets to the secondary objective components, source-coverage settings, and learning backends?
- **RQ4: Score randomization.** What predictive and security costs arise when we protect selected score blocks with stronger Laplace randomization?
- **RQ5: Scalability and runtime.** Does the primary solver match exhaustive search where enumeration is possible and remain practical as the candidate pool and number of target classes grow, and what communication reduction do we obtain at the evaluated scale?

A. Experimental Setup

Datasets. We use four multiclass cybersecurity datasets: CIC DDoS2019 (DDoS) [25], CICMalDroid2020 (MalDroid) [26], CIC Darknet2020 (Darknet) [27], and CIRA CIC DoHBrw 2020 (DoH) [28]. For each task, we construct 66 source contributors and let each contributor contribute one pretrained candidate detector. Keeping the candidate count fixed across datasets helps us separate dataset effects from changes in pool size.

We require every candidate detector to be multiclass, so each contributor receives at least two source classes and between 74 and 30,517 fitting records depending on the dataset and partition. Our main operating point uses a 50% selection-cost budget. If all candidate costs were identical, this budget would allow roughly 33 of the 66 detectors; in practice, the selected count depends on the logical byte cost b_i of each candidate. In RQ5, we separately vary the candidate pool from 16 to 256 to isolate solver scaling. Table II summarizes the four tasks.

Among the 13 classes observed in the reserved DDoS test set, one attack class is absent from all contributor-training partitions. We still include that class in macro F1 because it appears at test time, but we exclude it from the minimum evaluated class recall because no source detector could have learned it from the source data.

Protocol. For each dataset, we split the available training pool into disjoint 60%, 10%, and 30% partitions for source detector fitting, selection calibration, and shared-stacker fitting, respectively. After we fit the source detectors, the 60% source partitions are no longer used by the integration procedure. This instantiates the intended setting: the pretrained source detectors and their coarse metadata remain available, while the raw source records are no longer revisited.

The 10% calibration partition is used only to evaluate the pretrained candidate pool and select $\hat{S}$, whereas the 30% stacker-fitting partition is used only after $\hat{S}$ has been fixed. The test set remains sealed until both steps are complete. Unless stated otherwise, we repeat the complete protocol over five data splits, use LightGBM [29] for both source detectors and the shared stacker, set the selection-cost budget to 50%, use $\tau_L = 1$ and $n_{\min} = 2$, use inverse-square-root source-class weights, and assign equal weights to the three secondary components in Φ .

Baselines and reference point. We compare SILOSTITCH with two budget-matched selection rules and with the full detector pool. The first baseline is random selection. Within each data split, we generate 20 fixed random candidate orders and scan each order once, adding a candidate whenever its cost still fits the remaining budget. We average these 20 selections within the split. This baseline tells us whether an arbitrary budget-feasible subset is already sufficient.

The second baseline ranks candidates by mean calibrated Brier utility per selection-cost byte and greedily packs them in that order. This rule favors deployment-domain average utility but does not explicitly protect source class coverage. We use it to test the main reason for our ordered objective: a utility-driven selector may gain deployment-domain predictive quality by discarding unique source coverage. Finally, we use all 66 detectors as the unconstrained reference corresponding to the 100% selection-cost point.

Metrics and statistics. We use macro F1 as the main detection metric because every class contributes equally. This is important for imbalanced security traffic, where a frequent class can otherwise hide weak performance on a rare attack class. We also report accuracy, balanced accuracy, and minimum evaluated class recall. The last metric is the smallest recall

Table II

SCALE OF THE FOUR EVALUATION TASKS. EACH TASK USES 66 CONTRIBUTORS AND ONE PRETRAINED CANDIDATE DETECTOR PER CONTRIBUTOR. EVALUATED CLASSES COUNT LABELS IN THE SEALED TEST SET; DDoS CONTAINS ONE TEST CLASS ABSENT FROM SOURCE TRAINING.

Dataset	Training samples	Test samples	Features	Evaluated classes
DDoS	3,355,966	1,127,526	79	13
MalDroid	8,118	3,480	470	5
Darknet	99,066	42,459	76	11
DoH	1,005,708	431,034	29	4

among classes that both appear in the sealed test set and belong to the fixed source-training class set.

We report weighted source class coverage separately from these predictive metrics because it answers a different question: whether the selected set retains historical class support, not whether it predicts deployment records correctly. For the DDoS privacy experiment, we also report *attack rate among benign predictions*, defined as the fraction of predictions labeled benign that are actually attacks. We compute communication using Eqs. 6–8. Unless stated otherwise, we report means and 95% Student t confidence intervals over five repeated splits. When two conditions share a split, we pair their observations, and we apply Holm correction within each reported comparison family.

Experimental environment. We run the centralized replay on one Linux machine with Python 3.9.25, NumPy 1.26.4, scikit-learn 1.3.2, LightGBM 4.6.0, and XGBoost [30] 1.4.2. We record wall-clock runtime and the logical bytes exchanged by the evaluated protocol. Our centralized experiments replay the same logical workflow used in the analysis; in a distributed deployment, contributors compute and optionally randomize selected score blocks locally before transmission. The reported runtime therefore measures the centralized experimental replay, while the communication metric follows the logical contributor–coordinator workflow defined in the system model.

B. RQ1: Detection Quality under a Budget

We first ask whether the 50% selection-cost constraint provides a practically useful deployment point. A useful operating point should attain the primary source-coverage optimum under the budget, reduce communication substantially, and keep the final shared detector close to the full-pool reference. We therefore compare the methods at the 50% operating point and then inspect the full budget sweep in Fig. 2.

Detection quality and source coverage. At the 50% operating point, SILOSTITCH consumes 49.5% to 50.0% of the full-pool selection cost. Across all four tasks, its absolute mean macro-F1 gap from the full detector pool remains below 0.028. The gaps are at most 0.0052 on MalDroid, Darknet, and DoH, while DDoS has the largest gap at about 0.026. Thus, even the largest observed loss remains small relative to the roughly one-half reduction in selection cost.

The budget sweep in Fig. 2 shows that a larger selection-cost budget does not guarantee higher sealed-test macro F1 on every dataset. MalDroid improves most clearly as the budget

increases, DoH changes only slightly, and DDoS and Darknet are non-monotonic.

We next examine whether the communication savings degrade performance on the weakest class. Relative to the full pool, the 50% selected set changes accuracy by at most 0.0046 and minimum evaluated class recall by at most 0.0061 on MalDroid, Darknet, and DoH. On DDoS, the accuracy gap is 0.0296 but the minimum evaluated class-recall gap is only 0.0094. Darknet remains weak even under the full-pool reference, which reaches 0.0077 minimum evaluated class recall compared with 0.0058 after selection. Most importantly for our primary objective, SILOSTITCH attains the primary optimum W^* under the 50% budget on all four datasets.

Why coverage priority matters. The comparison with individual-utility ranking illustrates why deployment-domain average performance should not be optimized alone. On MalDroid and Darknet, SILOSTITCH has the higher mean macro F1, while the two methods are nearly identical on DoH. DDoS exposes the intended conflict most clearly: the utility-only baseline gains 0.0150 macro F1 over SILOSTITCH, but its weighted source class coverage falls to 0.910, whereas SILOSTITCH retains the maximum value of 1.0. In other words, the utility-only selector achieves a small deployment-domain predictive gain at the cost of reduced historical source class coverage. Our ordered objective explicitly rules out that exchange.

Overall, RQ1 shows the deployment advantage targeted by the paper: at roughly half the selection cost, we attain the primary source-coverage optimum and keep macro-F1 loss below 0.028 across all four tasks. The DDoS comparison also provides a concrete case in which the source-coverage priority is active rather than merely formal.

C. RQ2: Contributor Label Heterogeneity

The previous experiment starts from a fixed source candidate pool. We now study how the heterogeneity that creates such a pool can affect the integration problem before and after selection. We separate two effects: whether we can construct valid multiclass source detectors at all, and, when we can, whether stronger source-label imbalance consistently reduces the final macro F1.

We retain the original DDoS and Darknet contributor assignments for the original-layout analysis. For the controlled heterogeneity study, we generate Darknet and MalDroid partitions from Dirichlet distributions with $\alpha \in \{10, 1, 0.5, 0.1\}$, together

Table III

MACRO F1 AT THE 50% SELECTION-COST BUDGET. RANDOM SELECTION AND INDIVIDUAL-UTILITY RANKING RECEIVE THE SAME BUDGET; THE FULL DETECTOR POOL IS AN UNCONSTRAINED REFERENCE. ENTRIES REPORT MEANS AND 95% STUDENT t CONFIDENCE INTERVALS OVER FIVE REPEATED SPLITS.

Method	DDoS	MalDroid	Darknet	DoH
Full detector pool	0.6412 ± 0.0121	0.8449 ± 0.0123	0.5972 ± 0.0128	0.6073 ± 0.0042
Random selection	0.6079 ± 0.0287	0.8344 ± 0.0100	0.5931 ± 0.0071	0.6074 ± 0.0028
Individual utility ranking	0.6303 ± 0.0437	0.8374 ± 0.0132	0.5921 ± 0.0168	0.6069 ± 0.0042
SILOSTITCH	0.6153 ± 0.0453	0.8397 ± 0.0121	0.5953 ± 0.0148	0.6064 ± 0.0038




Figure 2. Macro F1 as a function of the selection-cost budget. We compare SILOSTITCH with the two budget-matched selection rules at 25%, 50%, and 75%; the star marks the unconstrained pool of 66 detectors. Error bars are 95% confidence intervals over five repeated splits after averaging the 20 random selections within each split. Each dataset uses its own vertical scale.

with an IID allocation. Smaller α creates stronger label imbalance. For the controlled partitions, we require every accepted source contributor to contain at least two training classes so that its pretrained candidate detector remains multiclass. Figure 3 reports both partition feasibility and the resulting detection quality.

When heterogeneity limits candidate construction. Darknet remains feasible for all ten seeds at every evaluated heterogeneity level, even as the measured contributor-label divergence rises from 0.0008 under IID allocation to 0.3516 at $\alpha = 0.1$. MalDroid behaves differently because it has only five source classes spread across the same 66 contributors. All ten IID and all ten $\alpha = 10$ partitions are usable, but only five $\alpha = 1$ partitions satisfy the two-class requirement. At $\alpha = 0.5$ and 0.1, none of the ten fixed seeds produces a valid partition despite allowing up to 100 generation attempts per seed. For these extreme settings, the difficulty arises before selective integration: under the evaluated generation procedure, no valid multiclass candidate pool was obtained.

Quality after a valid pool exists. Conditional on a usable candidate pool, stronger heterogeneity does not produce a monotonic decline in macro F1. Darknet stays between 0.5895 and 0.5976 at the 50% budget, while MalDroid stays between 0.8328 and 0.8375 over the feasible range. These narrow descriptive ranges suggest that, in this controlled experiment, increasing source-label imbalance affects candidate-pool feasibility more clearly than it affects aggregate detection quality

once a valid pool has already been constructed.

The original contributor layouts show a related effect at the class level. As the budget increases from 25% to 75%, Darknet macro F1 rises from 0.6162 to 0.6290, while its minimum evaluated class recall rises more substantially from 0.0452 to 0.0964. For DDoS, macro F1 changes from 0.6039 to 0.6104 over the same budget range, while minimum class recall increases from 0.0568 to 0.0773. Thus, additional selection-cost budget can be more visible in the weakest observed class than in an aggregate score.

When the coverage priority is non-binding. We also remove only the explicit source-coverage priority while keeping the same secondary refinement. Across 30 comparisons on the original assignments and 225 comparisons on the controlled assignments, the two variants choose the same set and therefore obtain the same macro F1, minimum class recall, communication, and source coverage. This is the complementary case to DDoS in RQ1: when the secondary search already reaches the maximum available source coverage, the hard priority is non-binding; when a utility-oriented selector would trade coverage away, the priority changes the result and prevents that loss.

D. RQ3: Selection Components and Learners

After the primary coverage value has been fixed, SILOSTITCH still needs to choose among coverage-equivalent detector sets and then fit the shared stacker. RQ3 therefore separates design choices at these two levels. We first remove the three




Figure 3. Effect of source-contributor label heterogeneity at the 50% selection-cost budget. Smaller Dirichlet α indicates stronger imbalance. Panel (a) counts valid partitions among ten fixed seeds, panel (b) reports contributor-label Jensen–Shannon divergence, and panels (c)–(d) report macro F1 for Darknet and MalDroid. Error bars are 95% confidence intervals over valid seeds; $n = 10$ except for MalDroid at $\alpha = 1$, where $n = 5$.

Table IV
MACRO F1 AT THE 50% SELECTION-COST BUDGET FOR EACH SOURCE DETECTOR AND STACKER COMBINATION. ENTRIES ARE MEANS OVER FIVE SPLITS; BOLD MARKS THE LARGEST DESCRIPTIVE MEAN FOR EACH DATASET.

Dataset	XGB →XGB	LGBM →XGB	XGB →LGBM	LGBM →LGBM
DDoS	0.617	0.635	0.633	0.615
MalDroid	0.806	0.833	0.827	0.840
Darknet	0.578	0.585	0.592	0.595
DoH	0.592	0.591	0.607	0.606

secondary objective components one at a time, then vary the source-coverage settings, and finally compare the source detector and stacker backends.

Secondary objective components. Recall that U measures record-level utility, C_{pred} measures conservative classwise utility, and D_{rep} measures how well the selected set represents the prediction patterns of the eligible pool. In every ablation, we keep the exact primary source coverage fixed and remove only one of these secondary terms. At the 50% budget, removing U lowers DDoS macro F1 by 0.0190, while removing D_{rep} raises it by 0.0181. Across datasets and budgets, however, the direction of these changes is not consistent, and no single removal remains statistically significant after Holm correction. We therefore interpret these terms as tie-breakers among coverage-equivalent sets rather than as individually universal sources of accuracy gain.

Source-coverage settings. We next vary the minimum source count $n_{\min} \in \{1, 2, 5, 10\}$ and replace the inverse-square-root class weights with uniform weights. Across two fixed Darknet layouts and two learner combinations, these changes leave the selected set, attained source coverage, macro F1, and minimum evaluated class recall unchanged. The tested Darknet result is therefore insensitive to these values of $n_{\min}$ and to the two evaluated class-weighting rules.

Learning backends. Table IV shows that backend choice creates larger descriptive differences than the source-coverage settings above. On DDoS, LightGBM source detectors with an

XGBoost shared stacker give the largest mean. On MalDroid and Darknet, LightGBM gives the largest mean at both levels. On DoH, XGBoost source detectors with a LightGBM stacker lead by a small margin. After Holm correction, only the two matched DoH comparisons between shared stackers remain statistically significant, and in both cases the LightGBM shared stacker is better than its XGBoost counterpart. We therefore treat the source detector and stacker backends as deployment tuning choices.

Taken together, RQ3 supports the intended separation between the primary and secondary stages. The maximum source class coverage is stable across these tested secondary variations, while the exact predictive ranking among coverage-equivalent sets remains task dependent.

E. RQ4: Score Randomization

The previous experiments use unperturbed selected scores. We now enable the optional local privacy mechanism introduced in Eq. 25 and measure the utility cost of stronger randomization. We fix the selected detector set and vary the per-block privacy parameter ϵ_b ; smaller ϵ_b means stronger Laplace noise. The selected-set and full-pool configurations use the same per-block ϵ_b . Because they release different numbers of detector blocks, this comparison does not hold the composed per-record privacy parameter $s\epsilon_b$ fixed. Theorem V.1 already gives the formal privacy guarantee, so here we focus on what the perturbation does to detection behavior.

Aggregate detection quality. Without added noise, the 50% selected set and the full detector pool obtain macro F1 values of 0.6153 and 0.6412, respectively. At the strongest evaluated perturbation, $\epsilon_b = 1$, these values fall to 0.5105 and 0.5298. Relative to their own nonprivate settings, the selected set loses 0.1048 and the full pool loses 0.1114. The two configurations therefore show similar absolute degradation at the evaluated per-block noise scale.

Attack rate among benign predictions. Macro F1 alone understates the security impact of the strongest perturbation. For the 50% selected set, the attack rate among benign predictions rises from 0.055 without noise to 0.262 at $\epsilon_b = 5$




Figure 4. Effect of Laplace score randomization on DDoS. Panel (a) reports macro F1 and panel (b) reports the attack rate among benign predictions, i.e., the fraction of predictions labeled benign that are attacks. From left to right, the per-block privacy parameter decreases and noise becomes stronger. We compare SILOSTITCH at the 50% selection-cost budget with the full detector pool at the same per-block ϵ_b ; their composed per-record privacy parameters differ with the number of released blocks. Error bars are 95% confidence intervals over the same five splits.

and 0.583 at $\epsilon_b = 1$. For the full pool, it rises from 0.056 to 0.193 and then to 0.644. Thus, at $\epsilon_b = 1$, more than half of the predictions labeled benign are actually attacks under both configurations. This is operationally more severe than the aggregate macro-F1 curve alone suggests. We therefore report both aggregate detection quality and the attack rate among benign predictions rather than relying on one summary score.

F. RQ5: Scalability and Runtime

The primary dynamic program is exact whenever the full recurrence completes without truncation, but Theorem V.3 shows that its state space can grow exponentially with the number of target classes C . RQ5 therefore checks three practical questions: whether the implementation matches exhaustive search when enumeration is possible, how runtime changes as the candidate and class dimensions grow, and what reduction in total logical communication we obtain at the evaluated scale.

Exactness audit. We first compare the dynamic program with exhaustive search at $C = 8$ and $M \in \{8, 10, 12, 16\}$. Across ten repetitions at each candidate-pool size, the solver reaches exactly the same primary coverage as exhaustive search in all 40 comparisons. Its median runtime ranges from 37.70 to 118.72 ms, and every 95th percentile remains below 125 ms. This audit checks the implementation against an enumerable ground truth; the larger sweeps below characterize runtime at greater scale.

Scaling behavior. We next vary the two problem dimensions separately in Fig. 5. With the class count fixed at 12, increasing

Table V
SYSTEM OUTCOME AT THE 50% SELECTION-COST BUDGET. WE DEFINE $\Delta F1 = F1_{full} - F1_{selected}$.

Dataset	Full $\rightarrow$ selected (MiB)	Saved	$\Delta F1$
DDoS	7040.05 $\rightarrow$ 3489.07	50.4%	0.0258
MalDroid	662.66 $\rightarrow$ 335.94	49.3%	0.0052
Darknet	3436.78 $\rightarrow$ 1732.29	49.6%	0.0019
DoH	2942.76 $\rightarrow$ 1488.51	49.4%	0.0009

the candidate pool from 66 to 256 raises the median selection time from 1.99 to 33.73 s. With 66 candidates fixed, increasing the class count from 12 to 20 raises the median from 2.95 to 100.10 s. The second increase is much sharper, which matches the structure of the bitmask solver: the candidate count enters linearly in the transition bound, while the class count controls the 2^C mask space. Thus, within the evaluated range, runtime is more sensitive to the class count than to the candidate-pool sizes considered here.

Total logical communication. Table V connects the solver study back to the main deployment objective. Across the four tasks, the 50% selection-cost setting reduces total logical communication by 49.3%–50.4% relative to the full pool, while the macro-F1 loss stays below 0.028. The three smallest gaps are at or below 0.0052; DDoS again has the largest gap but still retains the maximum weighted source class coverage from RQ1. Thus, the cost reduction reported in the abstract is not only a selector-level budget reduction: after we include the modeled candidate-upload, model-delivery, selected-score, and label-upload terms, total logical communication is still approximately halved.

The total centralized replay time for one split, including every evaluated selection rule and budget, ranges from 0.78 to 50.09 minutes depending on the dataset. Stacker fitting dominates MalDroid and Darknet, while selection takes the largest share on DDoS. At the main scale of 66 candidates, the primary selector itself finishes in seconds. Within the evaluated ranges, growth in the class count is the more visible scaling constraint, while the 66-candidate setting used in the main experiments remains practical.

VII. DISCUSSION AND CONCLUSION

In summary, we study selective integration for the setting in which raw historical source records are no longer available but pretrained intrusion detectors and coarse source metadata remain authorized for reuse. With SILOSTITCH, we align heterogeneous outputs, maximize weighted source class coverage under a hard deployment budget, refine deployment-domain predictive behavior without reducing that coverage, and fit one shared stacker from the selected outputs. Across four datasets, the 50% selection-cost setting reduces total logical communication by approximately half while keeping macro-F1 loss below 0.028 relative to the full detector pool.

When coverage priority matters. We use a strict coverage-first objective because source class coverage and deployment-domain predictive quality answer different questions. The




Figure 5. Scaling of the primary selector. Panel (a) fixes the class count at 12 and increases the candidate pool, while panel (b) fixes 66 candidates and increases the class count. The vertical axis is logarithmic; points are medians and upper bars are 95th percentiles over ten repetitions.

DDoS utility-only baseline shows a case in which this priority changes the result: it gains 0.0150 macro F1 but reduces weighted source class coverage from 1.0 to 0.910. The heterogeneity ablation shows the complementary case in which the priority is non-binding because the secondary search already reaches maximum coverage. Together, these results clarify the role of the ordered objective: it does not force a different set when the deployment-domain predictive search already preserves all available coverage, but it prevents coverage loss when deployment-domain utility would otherwise trade it away.

Our theoretical guarantees follow the same separation. The dynamic program solves the primary coverage problem exactly whenever the full recurrence completes without truncation at the implementation cap, and the local refinement preserves that optimum. The simultaneous calibration theorem remains valid after adaptive detector selection for the observed calibration mixture, while the mixture-shift result additionally assumes stable contributor-class mean Brier risk and bounded mixture change.

Deployment scope and scaling. The reported savings are modeled logical bytes rather than measurements of network latency or protocol overhead. The centralized replay nevertheless lets us separate the computational stages under identical data splits. At 66 candidates, primary selection finishes in seconds, but runtime grows quickly with the target class count because the bitmask state space depends on 2^C . The solver is therefore most suitable when the number of target classes remains moderate; larger ontologies may require a different primary solver or an approximation that explicitly relaxes exactness.

Privacy scope and utility tradeoff. Our local privacy guarantee applies to randomized selected score uploads and introduces a measurable utility cost. The DDoS experiment shows why we examine this cost at more than one level: under strong perturbation, macro F1 degrades, while the attack rate among benign predictions shows that more than half of benign-labeled predictions are attacks. A deployment should therefore tune score privacy using both aggregate predictive quality and security-critical error behavior.

Overall, our results show that when historical training

records cannot be revisited, a fixed subset of existing intrusion detectors can still be integrated in a budget-aware way without treating historical source coverage and deployment-domain predictive quality as interchangeable. SILOSTITCH turns that requirement into an explicit selection order and achieves substantial communication savings while keeping deployment-domain quality, privacy, and solver-scale limitations explicit.

REFERENCES

- [1] H. B. McMahan, E. Moore, D. Ramage, S. Hampson, and B. Agüera y Arcas, “Communication-efficient learning of deep networks from decentralized data,” in *Proceedings of the 20th International Conference on Artificial Intelligence and Statistics*, ser. Proceedings of Machine Learning Research, vol. 54. PMLR, 2017, pp. 1273–1282. [Online]. Available: <https://proceedings.mlr.press/v54/mcmahan17a.html>
- [2] T. Li, A. K. Sahu, M. Zaheer, M. Sanjabi, A. Talwalkar, and V. Smith, “Federated optimization in heterogeneous networks,” in *Proceedings of Machine Learning and Systems*, I. Dhillon, D. Papailiopoulos, and V. Sze, Eds., vol. 2. MLSys, 2020, pp. 429–450. [Online]. Available: https://proceedings.mlsys.org/paper_files/paper/2020/hash/1f5fe83998a09396e6e6477d9475ba0c-Abstract.html
- [3] T. Li, S. Hu, A. Beirami, and V. Smith, “Ditto: Fair and robust federated learning through personalization,” in *Proceedings of the 38th International Conference on Machine Learning*, ser. Proceedings of Machine Learning Research, vol. 139. PMLR, 2021, pp. 6357–6368. [Online]. Available: <https://proceedings.mlr.press/v139/li21h.html>
- [4] H. Zhao, W. Du, F. Li, P. Li, and G. Liu, “Fedprompt: Communication-efficient and privacy-preserving prompt tuning in federated learning,” in *ICASSP 2023-2023 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*. IEEE, 2023, pp. 1–5.
- [5] C. S. Johnson, M. L. Badger, D. A. Waltermire, J. Snyder, and C. Skorupka, “Guide to cyber threat information sharing,” National Institute of Standards and Technology, Tech. Rep. NIST Special Publication 800-150, 2016. [Online]. Available: <https://doi.org/10.6028/NIST.SP.800-150>
- [6] J. Xu, C. Li, Y. Zheng, and Z. Li, “ENTENTE: Cross-silo intrusion detection on network log graphs with federated learning,” in *Network and Distributed System Security Symposium*, 2026. [Online]. Available: <https://www.ndss-symposium.org/wp-content/uploads/2026-s93-paper.pdf>
- [7] E. GDPR, “General data protection regulation (gdpr),” *Cit. on*, vol. 4, 2018.
- [8] J. K.-W. Lee, P. Sattigeri, and G. W. Wornell, “Learning new tricks from old dogs: Multi-source transfer learning from pre-trained networks,” in *Advances in Neural Information Processing Systems*, vol. 32, 2019, pp. 4372–4382. [Online]. Available: <https://papers.nips.cc/paper/8688-learning-new-tricks-from-old-dogs-multi-source-transfer-learning-from-pre-trained-networks>
- [9] Q. Dong, A. Muhammad, F. Zhou, C. Xie, T. Hu, Y. Yang, S.-H. Bae, and Z. Li, “ZooD: Exploiting model zoo for out-of-distribution generalization,” in *Advances in Neural Information Processing Systems*, vol. 35, 2022. [Online]. Available: <https://proceedings.neurips>

- cc/paper_files/paper/2022/hash/cd305fdee96836d5cc1de94577d71b61-Abstract-Conference.html
- [10] X. Zhao, G. Sun, R. Cai, Y. Zhou, P. Li, P. Wang, B. Tan, Y. He, L. Chen, Y. Liang, B. Chen, B. Yuan, H. Wang, A. Li, Z. Wang, and T. Chen, "Model-GLUE: Democratized LLM scaling for a large model zoo in the wild," in *Advances in Neural Information Processing Systems*, vol. 37, 2024. [Online]. Available: <https://openreview.net/forum?id=M5JW7O9vc7>
 - [11] Y. Birman, S. Hindi, G. Katz, and A. Shabtai, "Cost-effective ensemble models selection using deep reinforcement learning," *Information Fusion*, vol. 77, pp. 133–148, 2022.
 - [12] L. Yang, W. Guo, Q. Hao, A. Ciptadi, A. Ahmadzadeh, X. Xing, and G. Wang, "CADE: Detecting and explaining concept drift samples for security applications," in *30th USENIX Security Symposium (USENIX Security 21)*. USENIX Association, 2021, pp. 2327–2344. [Online]. Available: <https://www.usenix.org/conference/usenixsecurity21/presentation/yang-limin>
 - [13] D. H. Wolpert, "Stacked generalization," *Neural Networks*, vol. 5, no. 2, pp. 241–259, 1992.
 - [14] Information Commissioner's Office, "Guidance on AI and data protection," Mar. 2023, version 2.0.17. [Online]. Available: <https://ico.org.uk/media/for-organisations/uk-gdpr-guidance-and-resources/artificial-intelligence/guidance-on-ai-and-data-protection-2-0.pdf>
 - [15] S. Thirumuruganathan, M. Nabeel, E. Choo, I. Khalil, and T. Yu, "SIRAJ: A unified framework for aggregation of malicious entity detectors," in *43rd IEEE Symposium on Security and Privacy*. IEEE, 2022, pp. 507–521.
 - [16] T. Qi, F. Wu, C. Wu, L. He, Y. Huang, and X. Xie, "Differentially private knowledge transfer for federated learning," *Nature Communications*, vol. 14, p. 3785, 2023. [Online]. Available: <https://www.nature.com/articles/s41467-023-38794-x>
 - [17] W. Qiu, Y. Feng, Y. Li, Y. Chang, K. Qian, B. Hu, Y. Yamamoto, and B. W. Schuller, "Fed-MStacking: Heterogeneous federated learning with stacking misaligned labels for abnormal heart sound detection," *IEEE Journal of Biomedical and Health Informatics*, vol. 28, no. 9, pp. 5055–5066, Sep. 2024.
 - [18] V. Pourahmadi, H. A. Alameddine, M. A. Salahuddin, and R. Boutaba, "Spotting anomalies at the edge: Outlier exposure-based cross-silo federated learning for DDoS detection," *IEEE Transactions on Dependable and Secure Computing*, vol. 20, no. 5, pp. 4002–4015, 2023.
 - [19] Y. Qing, Q. Yin, X. Deng, Y. Chen, Z. Liu, K. Sun, K. Xu, J. Zhang, and Q. Li, "Low-quality training data only? a robust framework for detecting encrypted malicious network traffic," in *Network and Distributed System Security Symposium*, 2024. [Online]. Available: <https://www.ndss-symposium.org/wp-content/uploads/2024-81-paper.pdf>
 - [20] R. Xie, J. Cao, E. Dong, M. Xu, K. Sun, Q. Li, L. Shen, and M. Zhang, "Rosetta: Enabling robust TLS encrypted traffic classification in diverse network environments with TCP-Aware traffic augmentation," in *32nd USENIX Security Symposium (USENIX Security 23)*. Anaheim, CA: USENIX Association, 2023, pp. 625–642. [Online]. Available: <https://www.usenix.org/conference/usenixsecurity23/presentation/xie>
 - [21] S. Yang, X. Zheng, J. Li, J. Xu, and E. C. H. Ngai, "CoLD: Collaborative label denoising framework for network intrusion detection," in *Network and Distributed System Security Symposium*, 2026. [Online]. Available: <https://www.ndss-symposium.org/wp-content/uploads/2026-s1950-paper.pdf>
 - [22] S. Maddock, G. Cormode, T. Wang, C. Maple, and S. Jha, "Federated boosted decision trees with differential privacy," in *Proceedings of the 2022 ACM SIGSAC Conference on Computer and Communications Security*, ser. CCS '22. New York, NY, USA: Association for Computing Machinery, 2022, pp. 2249–2263. [Online]. Available: <https://doi.org/10.1145/3548606.3560687>
 - [23] J. C. Duchi, M. I. Jordan, and M. J. Wainwright, "Local privacy and statistical minimax rates," in *2013 IEEE 54th Annual Symposium on Foundations of Computer Science*, 2013, pp. 429–438.
 - [24] C. Dwork, F. McSherry, K. Nissim, and A. Smith, "Calibrating noise to sensitivity in private data analysis," in *Theory of cryptography conference*. Springer, 2006, pp. 265–284.
 - [25] I. Sharafaldin, A. H. Lashkari, S. Hakak, and A. A. Ghorbani, "Developing realistic distributed denial of service (ddos) attack dataset and taxonomy," in *2019 International Carnahan Conference on Security Technology (ICCST)*. IEEE, 2019, pp. 1–8.
 - [26] S. MahdaviFar, A. F. A. Kadir, R. Fatemi, D. Alhadidi, and A. A. Ghorbani, "Dynamic android malware category classification using semi-supervised deep learning," in *IEEE International Conference on Dependable, Autonomic and Secure Computing*. IEEE, 2020, pp. 515–522.
 - [27] A. Habibi Lashkari, G. Kaur, and A. Rahali, "Didarknet: A contemporary approach to detect and characterize the darknet traffic using deep image learning," in *2020 the 10th International Conference on Communication and Network Security*, 2020, pp. 1–13.
 - [28] M. MontazeriShatoori, L. Davidson, G. Kaur, and A. H. Lashkari, "Detection of doh tunnels using time-series classification of encrypted traffic," in *2020 IEEE International Conference on Dependable, Autonomic and Secure Computing*. IEEE, 2020, pp. 63–70.
 - [29] G. Ke, Q. Meng, T. Finley, T. Wang, W. Chen, W. Ma, Q. Ye, and T.-Y. Liu, "Lightgbm: A highly efficient gradient boosting decision tree," *Advances in neural information processing systems*, vol. 30, pp. 3146–3154, 2017.
 - [30] T. Chen and C. Guestrin, "Xgboost: A scalable tree boosting system," in *Proceedings of the 22nd acm sigkdd international conference on knowledge discovery and data mining*, 2016, pp. 785–794.